

Operating Systems Assignment 2 Question 2

Yatharth Taneja | 2019346

Process

1. Installing and compiling the kernel v5.9.1

- 1.1 wget <link> to download the file from kernel.org
- 1.2 Installing the downloaded file to /usr/src/ using tar -xvz <file name> -C/usr/src/
- 1.3 installing necessary files and libraries like gcc, bison flex and use
sudo apt-get update
sudo apt-get upgrade
- 1.4 Make .config file using make menu config
- 1.5 Compile and install the kernel and modules using
make -j4
make modules_install install
- 1.6 reboot

2. Writing the syscall

- 2.1. Make a directory task_info in /usr/src
- 2.2. Create a file sh_task_info.c in this directory
- 2.3. Make a Makefile for the same
This is to make sure file is compiled and object code is created
- 2.4. Go to the /usr/src directory and edit it's makefile to add /task_info to the second occurrence of core-y:
This is to tell the compiler that the source files of our new system call (sys_sh_task_struct()) are in present in the task_info directory.
- 2.5. include the syscall in system call table present in arch/x86/entry/syscalls/
- 2.6 include the syscall in system call header file
- 2.7 compile the kernel again.

Description of code

1. sh_task_info.c file

- 1.1. Here we will have a pid as input therefore we will make a **task_struct** and assign values to it using the **find_task_by_vpid()**
- 1.2. Then we will get the data and allocate it to a buffer data using **snprintf()** function
- 1.3. This data can be printed on kernel log using **printk()** function
- 1.4. To write to a file we will first save the current segment and do a file open using **flip_open()**
- 1.5. We have the file_path as an argument we will copy it to a buffer using **strcpy_from_user** and write using **kernel_write()** with this buffer and data, data size as arguments.
- 1.6. At every important point, error handling is done.

2. test.c file

- 2.1. It calls the defined syscall sh_task_info and assigns a return value to a variable.
- 2.2. The syscall returns 0 for a successful run and -1 for errors

- 2.3. If there are any errors, the error is printed using `strerror(errno)` since `errno` is automatically assigned.

The inputs User should give

```
yatharth@ubuntu:~/Desktop/Assignment2/Q1$ ps
PID TTY TIME CMD
1983 pts/0 00:00:00 bash
2780 pts/0 00:00:00 ps
yatharth@ubuntu:~/Desktop/Assignment2/Q1$ top
top - 14:03:12 up 1:05, 1 user, load average: 0.00, 0.00, 0.00
Tasks: 311 total, 1 running, 310 sleeping, 0 stopped, 0 zombie
%Cpu(s): 1.1 us, 1.0 sy, 0.0 ni, 97.9 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
MiB Mem : 3907.3 total, 224.0 free, 1458.2 used, 2225.0 buff/cache
MiB Swap: 2048.0 total, 2048.0 free, 0.0 used, 2176.5 avail Mem

  PID USER      PR  NI    VIRT    RES    SHR   S %CPU  MEM%   TIME+  COMMAND
 1411 yatharth  20   0 370188 84928 48652 S   0.0   2.1   0:25.44 Xorg
1576 yatharth  20   0 4135236 256756 96292 S   2.3   6.4   0:49.06 gnome-s+
2141 yatharth  20   0 4550868 150048 53280 S   2.3   3.8   0:07.69 code
1975 yatharth  20   0 814356 51528 38928 S   1.3   1.3   0:06.25 gnome-t+
2783 yatharth  20   0 12104 4044 3124 R   1.0   0.1   0:00.09 top
2609 yatharth  20   0 1244872 83156 45328 S   0.7   2.1   0:09.14 sublim+
1697 yatharth  20   0 422664 31172 28792 S   0.3   0.8   0:00.52 gsd-pow+
2184 yatharth  20   0 14.5g 201368 103464 S   0.3   5.0   0:23.24 code
2603 root      20   0 0 0 0 I   0.3   0.0   0:02.89 kworker+
1 root      20   0 167796 11784 8560 S   0.0   0.3   0:25.96 systemd
2 root      20   0 0 0 0 S   0.0   0.0   0:00.08 kthreadd
3 root      0 -20 0 0 0 I   0.0   0.0   0:00.00 rcu_gp
4 root      0 -20 0 0 0 I   0.0   0.0   0:00.00 rcu_par+
6 root      0 -20 0 0 0 I   0.0   0.0   0:00.00 kworker+
9 root      0 -20 0 0 0 I   0.0   0.0   0:00.00 mm_perc+
10 root     20   0 0 0 0 S   0.0   0.0   0:00.07 ksoftirq+
11 root     20   0 0 0 0 I   0.0   0.0   0:01.29 rcu_sch+
12 root     rt  0 0 0 0 S   0.0   0.0   0:00.02 ntgrat+
13 root     -51  0 0 0 0 S   0.0   0.0   0:00.00 idle_in+
14 root     20   0 0 0 0 S   0.0   0.0   0:00.00 cpuhp/0
15 root     20   0 0 0 0 S   0.0   0.0   0:00.00 cpuhp/1
16 root     -51  0 0 0 0 S   0.0   0.0   0:00.00 idle_in+
```

The user should be able to run the program by using `make run`.

The user should give the pid and filepath as input in `test.c` and depending on the task whether it is running or not we will get the output.

We can check the process running either by `ps` or `top` command

Output

```
./a.out
Input pid :66

Input path : /home/yatharth/Desktop/out.txt
Syscall Returned 1
error :Invalid argument
root@ubuntu:/home/yatharth/Desktop/Assignment2/Q2# make run
gcc test.c
test.c: In function 'main':
test.c:18:9: warning: format '%s' expects argument of type 'char *', but argu
t 2 has type 'char (*)[256]' [-Wformat=]
18 | scanf("%s",&filepath);
   | ~~~~~^
   | |
   | | char (*)[256]
   | | char *
./a.out
Input pid :1

Input path : /home/yatharth/desktop/out.txt
Syscall Returned 1
error :Permission denied
root@ubuntu:/home/yatharth/Desktop/Assignment2/Q2#
```

If the pid is correct and syscall returns 0 and success message is printed

Else if there is an error sys call returns 1 along with the error type

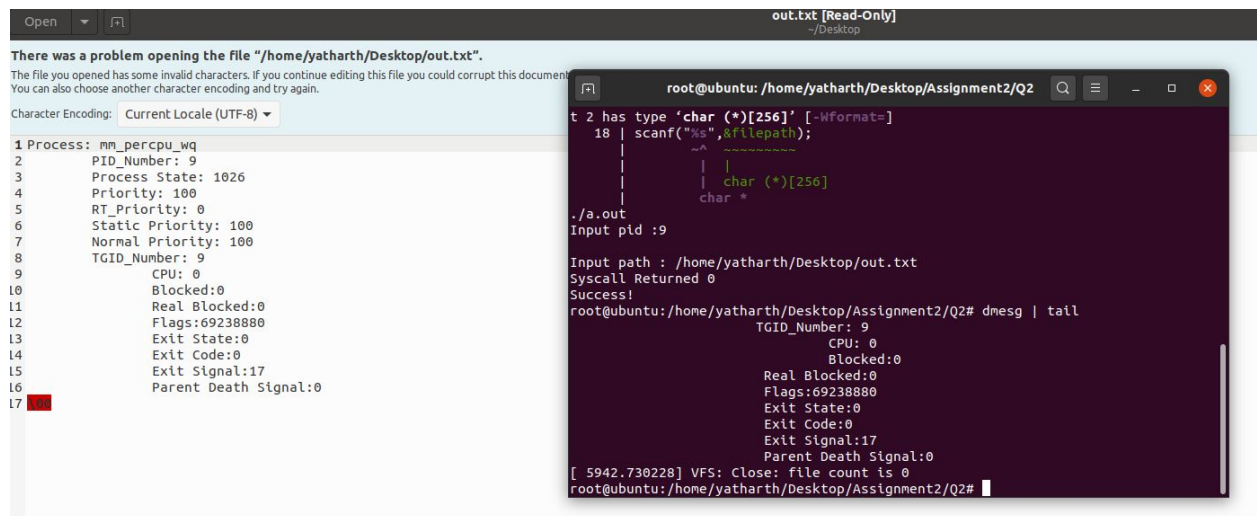
Here Invalid Argument is shown since 66 pid does not exist

And Permission denied is printed since I sent a wrong address. (desktop instead of Desktop)

```
[ 5782.837379] PID not found
[ 5815.565654] Process: systemd
PID_Number: 1
Process_State: 0
Priority: 100
RT_Priority: 0
Static_Priority: 120
Normal_Priority: 120
TGID_Numbers: 1
CPU: 0
Blocked: 67117312
Real_Blocked: 0
Flags: 4194560
Exit_State: 0
Exit_Code: 0
Exit_Signal: 17
Parent_Death_Signal: 0
[ 5815.565673] Cannot Open File
[ 5942.730174] Process: mm_percpu_wd
PID_Number: 9
Process_State: 4020
Priority: 100
RT_Priority: 0
Static_Priority: 100
Normal_Priority: 100
TGID_Numbers: 4
CPU: 0
Blocked: 0
Real_Blocked: 0
Flags: 69238880
Exit_State: 0
Exit_Code: 0
Exit_Signal: 17
Parent_Death_Signal: 0
[ 5942.730228] VFS: Close: file count is 0
root@ubuntu:/home/yatharth/Desktop/Assignment2/Q2#
```

This is the Output in the kernel log.

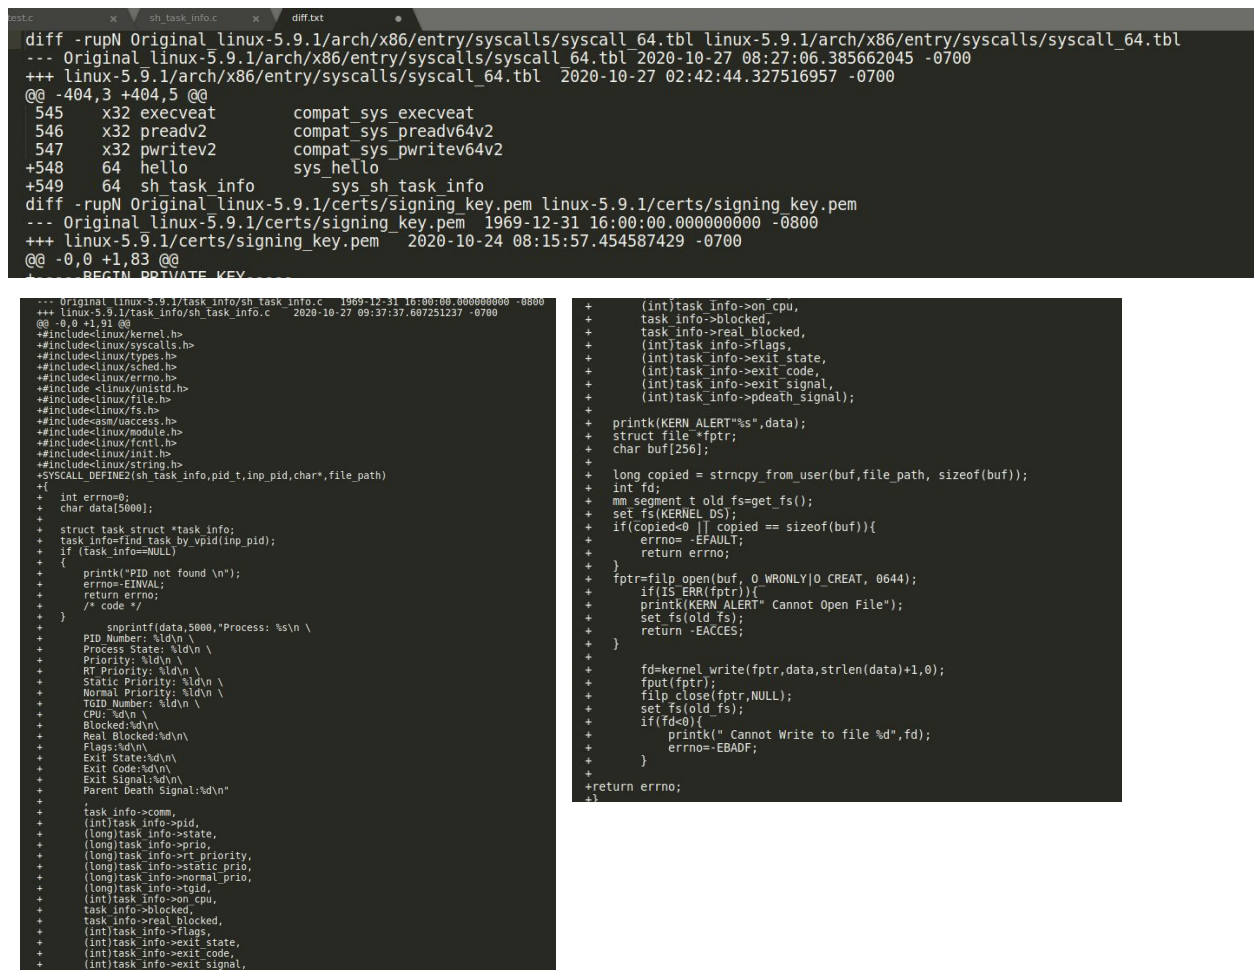
1. PID not found is for the Invalid argument
2. Cannot open is for the wrong address sent, (task_struct info is printed)
3. Log of the file successfully executed and closed



This is the combined output for the input pid 9 , dmesg | tail is used to print the last lines and on the left is the text file generated.

Note: You have to be in sudo -s mode for the file to be created

diff.txt These are the important changes shown in the patch file.



Error Handling

Incorrect pid	EINVAL	22	/* Invalid argument */
Unable to copy	EFAULT	14	/* Bad address */
Unable to open file	EACCES	13	/* Permission denied */
Unable to write to file	EBADF	9	/* Bad file number */

These are the four cases where errors are handled and you can see it in the syscall definition in the patch file the message inside `/* */` is the message that will be printed on the terminal as explained above in the description of code using `strerror(errno)`.